

CutMix: Regularization Strategy to Train Strong Classifiers with Localizable Features

Abstract

cnn

less discriminative parts of objects -> generalize better / better object localization capabilities.

regional dropout bad because -> remove informative pixels / information loss / inefficiency during training.

CutMix -> patches are cut and pasted among training images where the ground truth labels are also mixed proportionally to the area of the patches.

efficient use of training pixels / regularization effect of regional dropout

when used as a pretrained model, results in consistent performance

improves the model robustness against input corruptions and its out-of-distribution detection performances.

Introduction

기존 - including data augmentation / regularization techniques

to prevent a CNN from focusing too much on a small set of intermediate activations or on a small region on input images, random feature removal regularizations have been proposed.

ex) dropout

attend to the entire object region -> feature removal strategies -> generalization / localization

기존 dropout은 zeroed-out or filled with random noise -> greatly reducing the proportion of informative pixels

data hungry

maximally utilize the deleted regions / better generalization and localization -> CutMix

replace the removed regions with a patch from another image (See Table 1).

truth labels are also mixed proportionally to the number of pixels of combined images.





there is no uninformative pixel during training, making training efficient, while retaining the advantages of regional dropout to attend to non-discriminative parts of objects.

enhance localization ability

CutMix shares similarity with Mixup

Mixup samples tend to be unnatural

CutMix overcomes the problem by replacing the image region with a patch from another training image.

	ResNet-50	Mixup [48]	Cutout [3]	CutMix
Image				
Label	Dog 1.0	Dog 0.5 Cat 0.5	Dog 1.0	Dog 0.6 Cat 0.4
ImageNet Cls (%)	76.3 (+0.0)	77.4 (+1.1)	77.1 (+0.8)	<b>78.6</b> (+2.3)
ImageNet Loc (%)	46.3 (+0.0)	45.8 (-0.5)	46.7 (+0.4)	<b>47.3</b> (+1.0)
Pascal VOC Det (mAP)	75.6 (+0.0)	73.9 (-1.7)	75.1 (-0.5)	<b>76.7</b> (+1.1)

CutMix also enhances the model robustness and alleviates the over-confidence issue [13, 22] of deep networks.

Related Works

Regional Dropout

Synthesizing Training Data

Mixup

Tricks for training deep networks

Methods such as weight decay,  
dropout [34], and Batch Normalization  
noises to the internal features of CNNs [17, 8, 46] or adding  
extra path to the architecture

CutMix

Algorithm

$$\tilde{x} = \mathbf{M} \odot x_A + (\mathbf{1} - \mathbf{M}) \odot x_B$$
$$\tilde{y} = \lambda y_A + (1 - \lambda) y_B,$$

simple

negligible computational overhead

Discussion

	Mixup	Cutout	CutMix
Usage of full image region	✓	✗	✓
Regional dropout	✗	✓	✓
Mixed image & label	✓	✗	✓

What does model learn with CutMix?

Cutout - lets a model focus on less discriminative parts of the object

CutMix - Cutout + localize

Analysis on validation error

stabilizing  
diverse training samples reduce overfitting.

Experiments

Robustness and Uncertainty				
	Baseline	Mixup	Cutout	CutMix
Top-1 Acc (%)	8.2	24.4	11.5	<b>31.0</b>

Method	TNR at TPR 95%	AUROC	Detection Acc.
Baseline	26.3 (+0)	87.3 (+0)	82.0 (+0)
Mixup	11.8 (-14.5)	49.3 (-38.0)	60.9 (-21.0)
Cutout	18.8 (-7.5)	68.7 (-18.6)	71.3 (-10.7)
CutMix	<b>69.0 (+42.7)</b>	<b>94.4 (+7.1)</b>	<b>89.1 (+7.1)</b>

Conclusion

strong classification and localization  
easy to implement / no computational overhead